

```
In [1]: #Aim:To perform and find the accuracy of Logistic Regression
```

```
In [ ]: #Name: Achal Subhash Kharwade
#Roll No: 36
#Sec: B
#Date :09-10-2023
```

```
In [4]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings('ignore')
import os
```

```
In [5]: os.getcwd()
```

Out[5]: 'C:\\Users\\Lenovo\\DSS 5th Sem'

```
In [29]: os.chdir("D:\\DSS\\DSS PRAC PG")
```

```
In [7]: df=pd.read_csv("framingham.csv")
```

```
In [8]: df.head()
```

Out[8]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate
0	1	39	4.0	0	0.0	0.0	0	0	0	195.0	106.0	70.0	26.97	80.1
1	0	46	2.0	0	0.0	0.0	0	0	0	250.0	121.0	81.0	28.73	95.1
2	1	48	1.0	1	20.0	0.0	0	0	0	245.0	127.5	80.0	25.34	75.1
3	0	61	3.0	1	30.0	0.0	0	1	0	225.0	150.0	95.0	28.58	65.1
4	0	46	3.0	1	23.0	0.0	0	0	0	285.0	130.0	84.0	23.10	85.1

```
In [9]: df.describe()
```

Out[9]:

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol
count	4240.000000	4240.000000	4135.000000	4240.000000	4211.000000	4187.000000	4240.000000	4240.000000	4240.000000	4190.000000
mean	0.429245	49.580189	1.979444	0.494104	9.005937	0.029615	0.005896	0.310613	0.025708	236.69952
std	0.495027	8.572942	1.019791	0.500024	11.922462	0.169544	0.076569	0.462799	0.158280	44.59128
min	0.000000	32.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	107.00000
25%	0.000000	42.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	206.00000
50%	0.000000	49.000000	2.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	234.00000
75%	1.000000	56.000000	3.000000	1.000000	20.000000	0.000000	0.000000	1.000000	0.000000	263.00000
max	1.000000	70.000000	4.000000	1.000000	70.000000	1.000000	1.000000	1.000000	1.000000	696.00000

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
#   Column              Non-Null Count  Dtype
---  -
0   male                 4240 non-null   int64
1   age                  4240 non-null   int64
2   education            4135 non-null   float64
3   currentSmoker        4240 non-null   int64
4   cigsPerDay           4211 non-null   float64
5   BPMeds               4187 non-null   float64
6   prevalentStroke       4240 non-null   int64
7   prevalentHyp         4240 non-null   int64
8   diabetes             4240 non-null   int64
```

```

9   totChol      4190 non-null float64
10  sysBP        4240 non-null float64
11  diaBP        4240 non-null float64
12  BMI          4221 non-null float64
13  heartRate    4239 non-null float64
14  glucose      3852 non-null float64
15  TenYearCHD   4240 non-null int64
dtypes: float64(9), int64(7)
memory usage: 530.1 KB

```

```
In [11]: df.isna().sum()
```

```

Out[11]: male          0
age            0
education      105
currentSmoker  0
cigsPerDay     29
BPMeds         53
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        50
sysBP          0
diaBP          0
BMI            19
heartRate      1
glucose        388
TenYearCHD     0
dtype: int64

```

```
In [12]: df
```

```

Out[12]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate
0	1	39	4.0	0	0.0	0.0	0	0	0	195.0	106.0	70.0	26.97	
1	0	46	2.0	0	0.0	0.0	0	0	0	250.0	121.0	81.0	28.73	
2	1	48	1.0	1	20.0	0.0	0	0	0	245.0	127.5	80.0	25.34	
3	0	61	3.0	1	30.0	0.0	0	1	0	225.0	150.0	95.0	28.58	
4	0	46	3.0	1	23.0	0.0	0	0	0	285.0	130.0	84.0	23.10	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
4235	0	48	2.0	1	20.0	NaN	0	0	0	248.0	131.0	72.0	22.00	
4236	0	44	1.0	1	15.0	0.0	0	0	0	210.0	126.5	87.0	19.16	
4237	0	52	2.0	0	0.0	0.0	0	0	0	269.0	133.5	83.0	21.47	
4238	1	40	3.0	0	0.0	0.0	0	1	0	185.0	141.0	98.0	25.60	
4239	0	39	3.0	1	30.0	0.0	0	0	0	196.0	133.0	86.0	20.91	

4240 rows × 16 columns

```
In [13]: # Missing value Treatment
# Since, 'glucose' and 'education' columns hadd a significant amount of null values , so we replaced them with the
```

```
In [14]: df['glucose'].fillna(value = df['glucose'].mean(),inplace=True)
```

```
In [15]: df['education'].fillna(value = df['education'].mean(),inplace=True)
```

```
In [16]: df['heartRate'].fillna(value = df['heartRate'].mean(),inplace=True)
```

```
In [17]: df['BMI'].fillna(value = df['BMI'].mean(),inplace=True)
```

```
In [18]: df['cigsPerDay'].fillna(value = df['cigsPerDay'].mean(),inplace=True)
```

```
In [19]: df['totChol'].fillna(value = df['totChol'].mean(),inplace=True)
```

```
In [20]:
```

```
In [20]: df['BPMeds'].fillna(value = df['BPMeds'].mean(),inplace=True)
```

```
In [21]: df.isna().sum()
```

```
Out[21]: male          0
age          0
education    0
currentSmoker 0
cigsPerDay   0
BPMeds       0
prevalentStroke 0
prevalentHyp 0
diabetes     0
totChol      0
sysBP        0
diaBP        0
BMI          0
heartRate    0
glucose      0
TenYearCHD   0
dtype: int64
```


```
In [22]: #Spilting the dependent and independent variables
x = df.drop("TenYearCHD",axis=1)
y=df['TenYearCHD']
```

```
In [23]: x #checking the features
```

```
Out[23]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate
0	1	39	4.0	0	0.0	0.000000	0	0	0	195.0	106.0	70.0	26.97	
1	0	46	2.0	0	0.0	0.000000	0	0	0	250.0	121.0	81.0	28.73	
2	1	48	1.0	1	20.0	0.000000	0	0	0	245.0	127.5	80.0	25.34	
3	0	61	3.0	1	30.0	0.000000	0	1	0	225.0	150.0	95.0	28.58	
4	0	46	3.0	1	23.0	0.000000	0	0	0	285.0	130.0	84.0	23.10	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
4235	0	48	2.0	1	20.0	0.029615	0	0	0	248.0	131.0	72.0	22.00	
4236	0	44	1.0	1	15.0	0.000000	0	0	0	210.0	126.5	87.0	19.16	
4237	0	52	2.0	0	0.0	0.000000	0	0	0	269.0	133.5	83.0	21.47	
4238	1	40	3.0	0	0.0	0.000000	0	1	0	185.0	141.0	98.0	25.60	
4239	0	39	3.0	1	30.0	0.000000	0	0	0	196.0	133.0	86.0	20.91	

4240 rows × 15 columns



```
In [24]: #Train Test Split
```

```
In [25]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=42)
```

```
In [26]: y_train
```

```
Out[26]: 1427    0
3257    0
3822    0
1263    0
3575    0
..
3444    0
466     0
3092    0
3772    0
860     0
Name: TenYearCHD, Length: 3392, dtype: int64
```

```
In [27]: #Logistic Regression Algorithm
```

```
In [28]: from sklearn.linear model import LogisticRegression
```

```
from sklearn.linear_model import LogisticRegression
model=LogisticRegression().fit(x_train,y_train)
model.score(x_train,y_train)
```

Out[28]: 0.8484669811320755

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js